

Course 2 - Hyperparameters, Regularization and Optimization

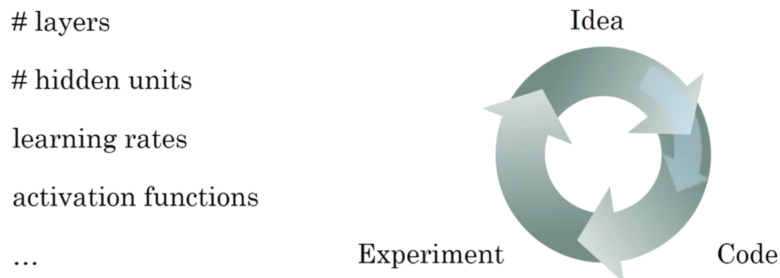
Mohammad Khalaji

August 14, 2020

1 Week 1

1.1 Train/Dev/Test Sets

We should always consider the fact that applied ML is an extremely iterative process with many parameters that affect our final outcome. It is virtually impossible to guess those parameters correctly in our first attempt.

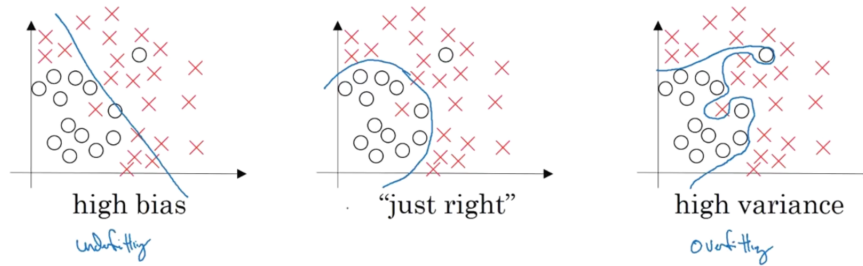


In the previous era of machine learning, it was common to divide all the data and split it according to a 70%/30% proportion, and assign them as your train and test sets respectively. Nowadays, however, with the emergence of big data, dev and test sets occupy a smaller percentage of our data. For example, when we have a dataset with 100000 examples, an 20% dev set is too big! As a result, the proportions nowadays look like 99.5%/0.25%/0.25%.

Not having a test set might be okay (only train/dev sets).

1.2 Bias/Variance

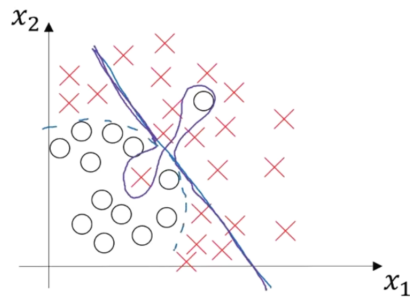
The simplest explanation:



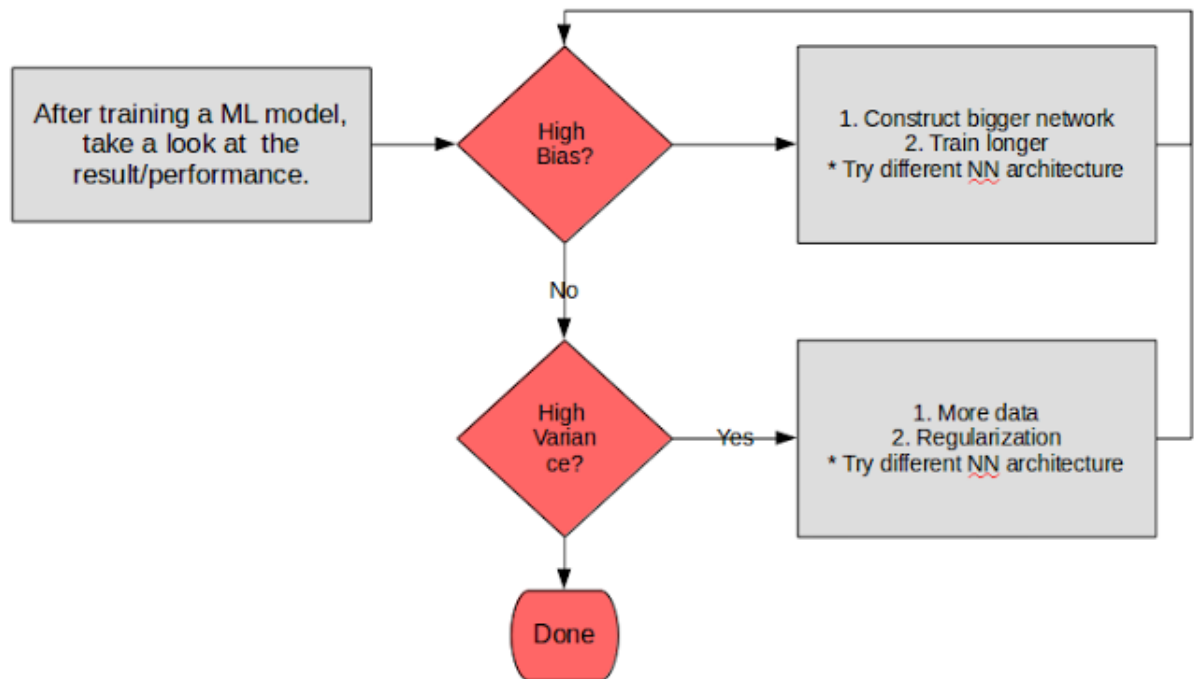
Consider another example for a cat classifier.

	High Variance	High Bias	High Bias High Variance	Low Bias Low Variance
Train Error	1%	15%	15%	0.5%
Test Error	11%	16%	30%	1.5%

The following classifier has both high bias and high variance. High bias because it is a mostly linear classifier, and it mostly underfits; high variance because in some occasions it shows overfitting behavior.



1.3 Basic Recipe for Machine Learning



1.4 Regularization

If you suspect that your neural network has a high variance problem, meaning that it overfits your data, regularization is one of the first things you should try. The other option is to get more training data, but it is not always viable.

Remember that in logistic regression, we were trying to minimize $J(w, b)$. Now we add the regularization term to the previous J .

$$J(w, b) = \frac{1}{m} \sum_{i=1}^m L(\hat{y}^{(i)}, y^{(i)}) + \frac{\lambda}{2m} \|w\|_2^2$$

It's called L_2 regularization because it uses L_2 norm, which is:

$$\|w\|_2^2 = \sum_{j=1}^{n_x} w_j^2 = w^T w$$

It is also possible that we use L_1 regularization with the regularization term:

$$\frac{\lambda}{2m} \sum_{j=1}^{n_x} |w_j| = \frac{\lambda}{2m} \|w\|_1$$

Note that also a regularization term for bias is possible, but since b is not as high-dimensional as w , its regularization term, which is equal to $\frac{\lambda}{m} b^2$, is usually omitted.

In a more complex neural network, though, a more generalized equation will be like this:

$$J(W^{[1]}, b^{[1]}, \dots, W^{[L]}, b^{[L]}) = \frac{1}{m} \sum_{i=1}^m L(\hat{y}^{(i)}, y^{(i)}) + \frac{\lambda}{2m} \sum_{l=1}^L \|W^{[l]}\|_F^2$$

$\|W^{[l]}\|_F^2$ is called the *Frobenius Norm* of the matrix $W^{[l]}$, and knowing that the shape of $W^{[l]}$ is $(n^{[l-1]}, n^{[l]})$ it's computed like this:

$$\|W^{[l]}\|_F^2 = \sum_{i=1}^{n^{[l-1]}} \sum_{j=1}^{n^{[l]}} (w_{ij}^{[l]})^2$$

Regularization is also applied to the gradients (it is called *weight decay*):

$$dw^{[l]} = (\text{from backprop}) + \frac{\lambda}{m} w^{[l]}$$

1.5 Why Regularization Reduces Overfitting

Consider a situation where you have a neural network that overfits your data. Remember the original non-regularized cost function:


$$J = \frac{1}{m} \sum_{i=1}^m L(\hat{y}^{(i)}, y^{(i)})$$

In order to reduce overfitting, we penalized J with a regularization term:

$$J = \frac{1}{m} \sum_{i=1}^m L(\hat{y}^{(i)}, y^{(i)}) + \frac{\lambda}{2m} \sum_{l=1}^L \|W^{[l]}\|_F^2$$


Intuition: If we set λ to be a really big number, our minimization algorithm will seek to set $W^{[l]} = 0$, which essentially means that a lot of hidden units will be zeroed out, which means that their impact will be reduced, and we will end up with a simpler neural network that is closer to the "High Bias" realm than it is to the "High Variance" realm.

Another Intuition: Consider using hyperbolic tangent as the activation function. If we regularize our weights, they will become smaller, and as a result, $z = wx + b$ will be smaller. With smaller zs , we will be wandering in the "linear section" of \tanh , which means that each layer will be *roughly* linear, which leaves our network having higher bias.



1.6 Dropout Regularization

For each node in the network, set some probability for eliminating it. After selecting the to-be-eliminated nodes, remove all the outgoing and incoming links from and to that node, and then do backpropagation on the diminished network.



Since you are training a much smaller network *on each example*, you will end up with a regularized network.

Dropout Implementation (Inverted Dropout): Keep in mind that this works for both the vectorized version and the non-vectorized version. Given $l = 3$, $keepProb = 0.8$:

$$d^{[3]} = np.random.rand(a^{[3]}.shape)$$

$$a^{[3]} = d^{[3]} * a^{[3]}$$

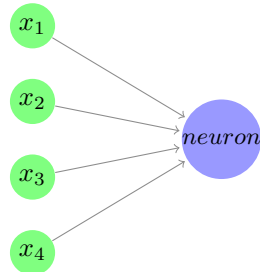
$$a^{[3]} /= keepProb$$

What does the last line do? It keeps the expected value of $a^{[3]}$ stable so as to not reduce the value of $z^{[4]} = W^{[4]}a^{[3]} + b^{[4]}$.

Obviously, there will be no dropout in the test phase, because we do not want to use a random element in our final computation.

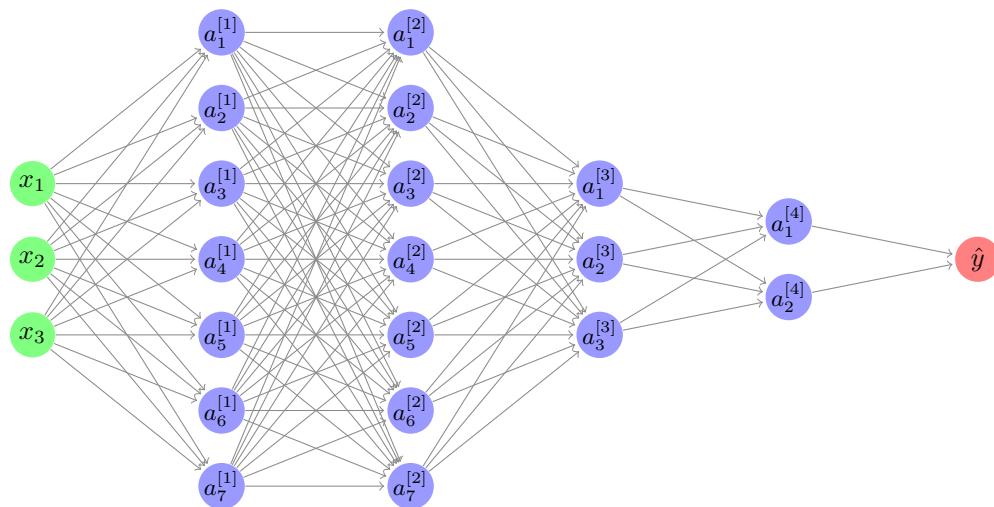
1.7 Understanding Dropout

Consider this one particular purple neuron:



As you can see, it relies on 4 features for its value, as a result, using dropout makes the network reluctant to rely heavily on *one* particular feature, employing a more *egalitarian* approach. Dropout has a similar effect to L_2 regularization's.

Another worthy note: It is good practice to keep *keepProb* variable according to the layer:

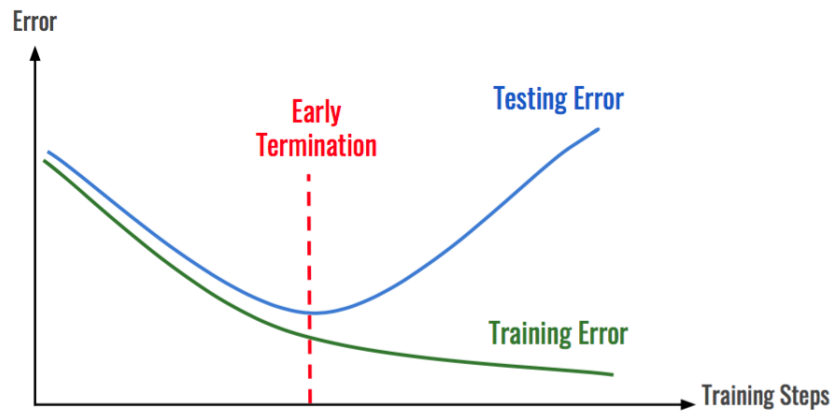


You might choose to use a lower *keepProb* for layers that have more parameters, because you are more worried about their overfitting ($l = 2$ for example).

1.8 Other Regularization Methods

Data Augmentation: Rotating, flipping, distorting and scaling train data images, for example.

Early Stopping:



It does have a downside though, according to the *orthogonalization* principle, we want to worry about one thing at a time. So worrying about the minimization of J , and worrying about not overfitting should be handled separately. Early stopping prevents us from doing that.

1.9 Normalizing Inputs

1. Subtract Mean:

$$\mu = \frac{1}{m} \sum_{i=1}^m X^{(i)}$$

$$X^{(i)} = X^{(i)} - \mu \quad \forall 1 \leq i \leq m$$


2. Normalize Variance:

$$\sigma^2 = \frac{1}{m} \sum_{i=1}^m X^{(i)} * * 2$$

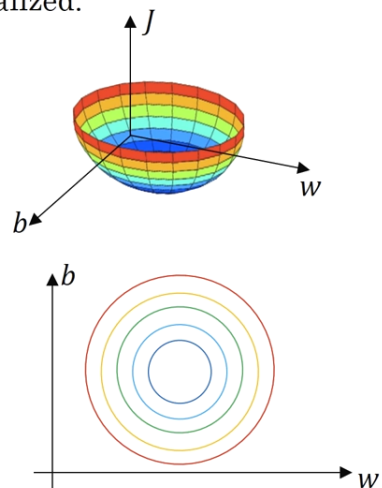
$$X^{(i)} = X^{(i)} / \sigma^2 \quad \forall 1 \leq i \leq m$$

Important: Use the same μ, σ^2 to normalize your test set. DO NOT normalize train/dev/test sets with different values of μ, σ^2 .

Unnormalized:




Normalized:



1.10 Exploding/Vanishing Gradients

Consider the following deep network, where $g^{[l]}(z) = z$, $b^{[l]} = 0$ in all layers:



$$\hat{y} = W^{[L]}W^{[L-1]} \dots W^{[2]}W^{[1]}X$$

If we have either of the scenarios below, the final \hat{y} will be either exponentially big or small, making it very difficult to train the network.

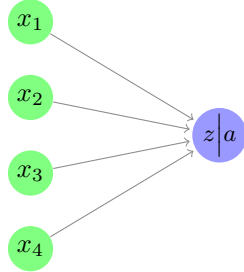
$$W^{[l]} = \begin{bmatrix} 1.5 & 0 \\ 0 & 1.5 \end{bmatrix} \Rightarrow \hat{y} = W^{[L]}(1.5)^{L-1}X$$

$$W^{[l]} = \begin{bmatrix} 0.5 & 0 \\ 0 & 0.5 \end{bmatrix} \Rightarrow \hat{y} = W^{[L]}(0.5)^{L-1}X$$

There is a *partial* solution to this problem, which is careful initialization of the weights.

1.11 Weight Initialization in a Deep Network

Consider this single neuron example, $b = 0$:



$$z = w_1x_1 + \cdots + w_nx_n$$

Now we don't want z to become too big or too small, so with larger n , we want smaller w_i s. One solution is to set $\text{Variance}(w_i) = \frac{1}{n}$.

$$w^{[l]} = np.random.randn(shape) * np.sqrt(\frac{1}{n^{[l-1]}})$$

(For ReLU activation functions, $\text{Variance}(w_i) = \frac{2}{n^{[l-1]}}$ works better)

(For \tanh activation function, $\text{Variance}(w_i) = \sqrt{\frac{1}{n^{[l-1]}}}$ and $\text{Variance}(w_i) = \sqrt{\frac{2}{n^{[l-1]} + n^{[n]}}}$ work better)

1.12 Numerical Approximation of Gradients

$$g(\theta) \approx \frac{f(\theta + \epsilon) - f(\theta - \epsilon)}{2\epsilon}$$

1.13 Gradient Checking

To verify your backprop implementation is correct.

- Take all your parameters, $W^{[1]}, b^{[1]}, \dots, W^{[L]}, b^{[L]}$, and reshape them into a big vector θ .
- Take all your gradients, $dW^{[1]}, db^{[1]}, \dots, dW^{[L]}, db^{[L]}$, and reshape them into a big vector $d\theta$.
- Now the question is: Is $d\theta$ the gradient of $J(\theta)$?

Here's how it's done:

- For each i :

$$d\theta_{approx}[i] = \frac{J(\theta_1, \theta_2, \dots, \theta_i + \epsilon, \dots) - J(\theta_1, \theta_2, \dots, \theta_i - \epsilon, \dots)}{2\epsilon}$$
$$\approx d\theta[i] = \frac{\partial J}{\partial \theta_i}$$

- Now that you have the two vectors $d\theta$ and $d\theta_{approx}$, compute this (the denominator is just a normalizer):

$$\frac{\|d\theta_{approx} - d\theta\|_2}{\|d\theta_{approx}\|_2 + \|d\theta\|_2}, \epsilon = 10^{-7}$$

- If the result is 10^{-7} or smaller, your gradient is very likely to be working properly. However, you need to double check if the value is too big.